

# **KLASIFIKASI GENRE MUSIK**

## **KELOMPOK 1**

**ANGGOTA KELOMPOK:**

**KEANE CLEMENT (2602084981)**

**KEVIN MAXWELL (2602083524)**

**DERREN MALAKA (2602081582)**

**JOSEL LINUS (2602076216)**

**RAFEL SUPRATMAN (2602073675)**

**HENDRA PUTRA (2602085095)**

**KELSTEN WUISAN (2602080932)**

# Introduksi

Kami akan membuat model machine learning yang dapat mengklasifikasikan sebuah lagu berdasarkan genre, mulai dari pop hingga rock dengan menggunakan machine learning.

# Problem & Solution

## Problem

- Genre: a style, especially in the arts, that involves a particular set of characteristics.
- Orang-orang menggunakan genre untuk mengkategorikan lagu.
- Tidak ada metode matematis yang bisa secara langsung mengkategorikan sebuah lagu ke sebuah genre.

## Solution

- Menggunakan machine learning untuk secara otomatis mengkategorikan sebuah lagu ke sebuah genre.
- Menggunakan dataset berisikan instances lagu dalam bentuk waveform dan labelnya.
- Proses waveform lagu menggunakan librosa.

# Dataset

Dataset yang akan kami gunakan untuk melaksanakan model machine learning ini adalah GTZAN Dataset - Music Genre Classification yang kami dapatkan dari kaggle.

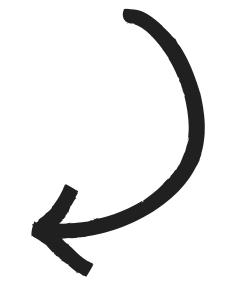
Link Dataset:

<https://www.kaggle.com/datasets/andradaolteanu/gtzan-dataset-music-genre-classification/code>

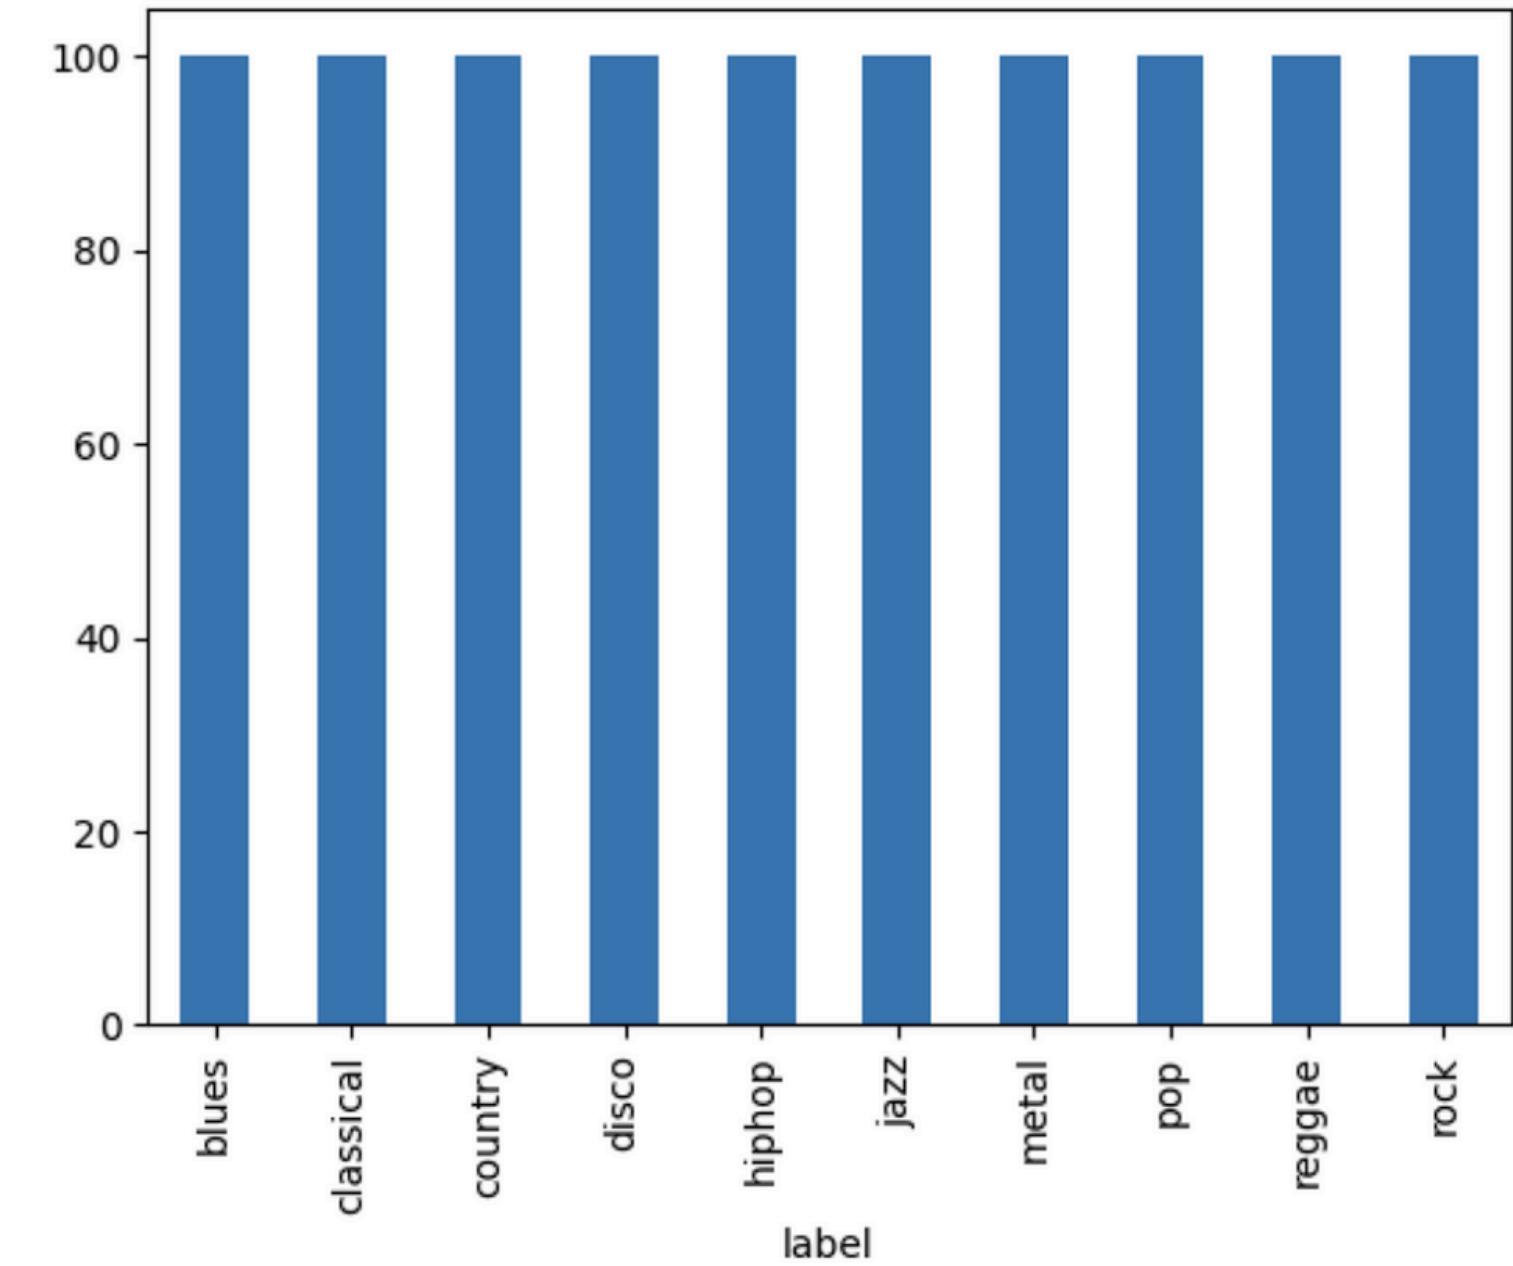
# Data Flow

Waveform -> Librosa

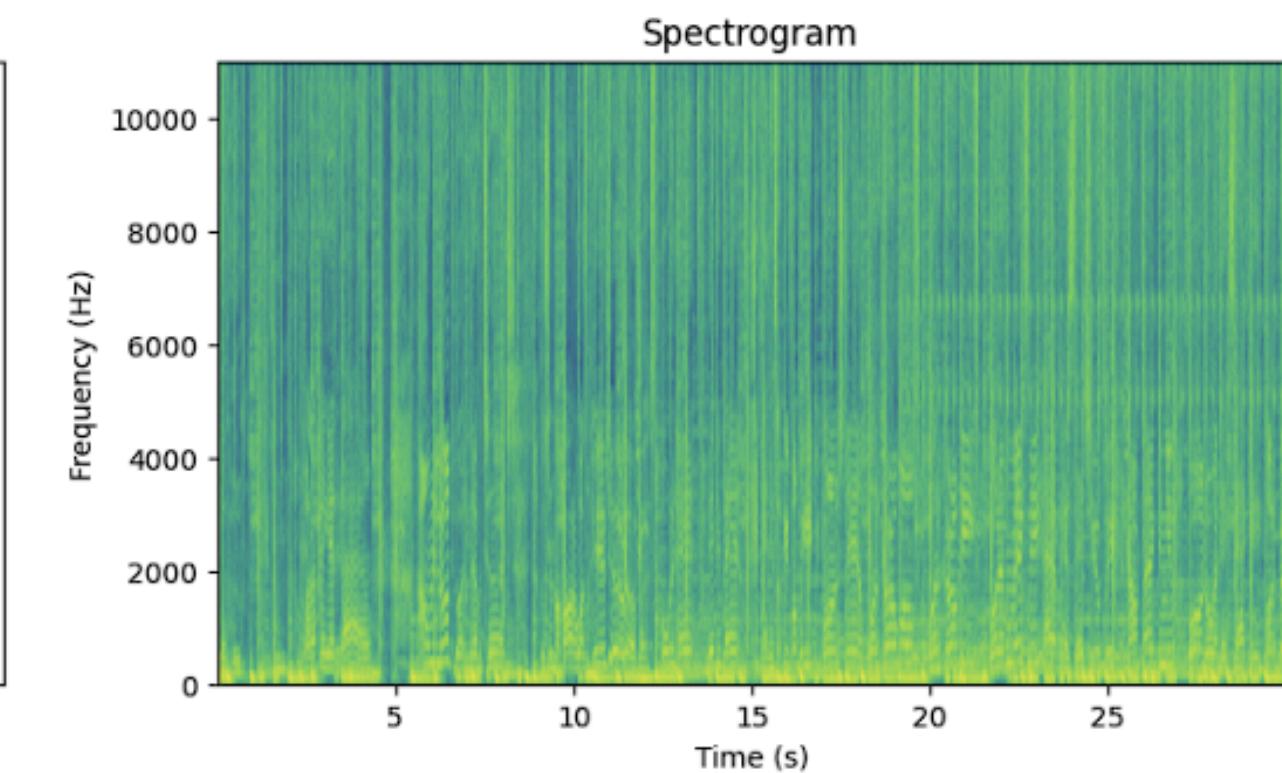
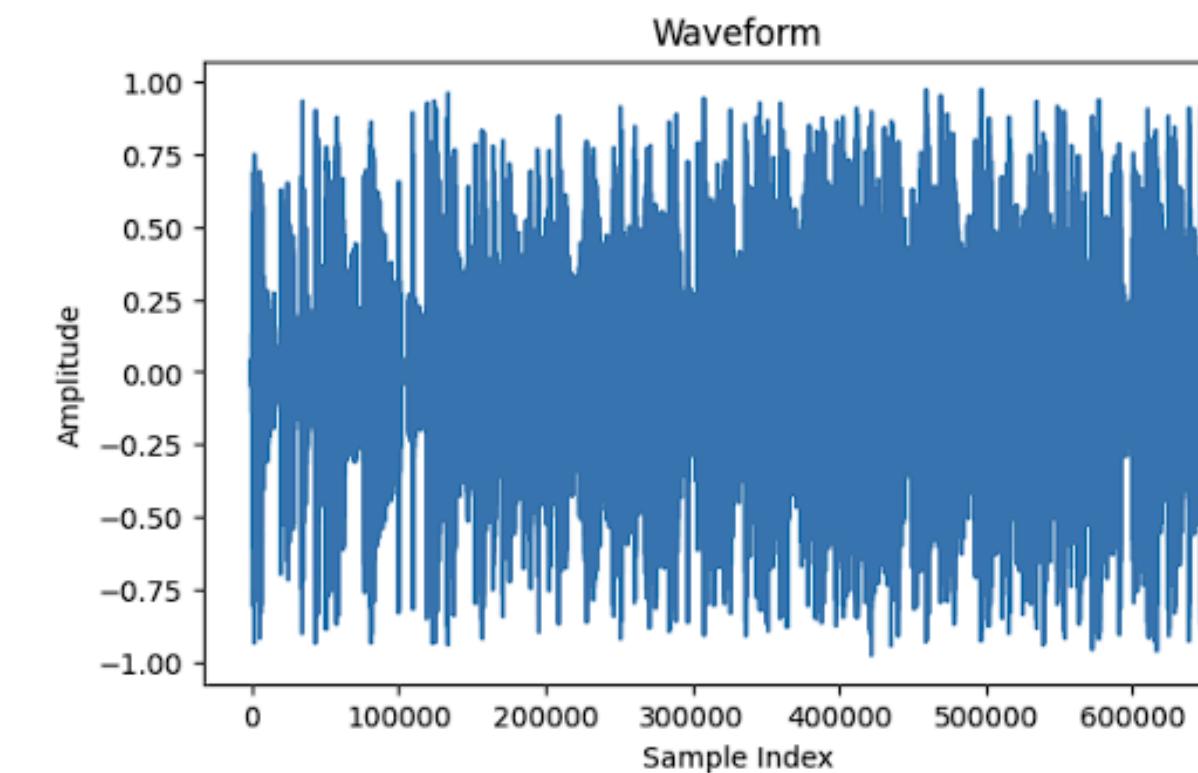
Classifier <- Audio Feature



# Data Exploration



# Data Exploration



# Feature Extraction

**Zero Crossing Rate**  
**Harmonic & Perceptual**  
**Spectral Centroid**  
**Spectral Bandwidth**  
**Spectral Rolloff**  
**MFCC**  
**Chroma**  
**Root Mean Square**  
**Tempo**

# Data Preprocessing

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder

from sklearn.model_selection import train_test_split

df = df_3sec
df = df.drop(columns = 'filename')
le = LabelEncoder()
df['label_encoded'] = le.fit_transform(df['label'])

for i, label in enumerate(df['label'].unique()):
    print(i, label, end=" | ")
df.pop('label')

scaler = MinMaxScaler()
X = df.drop(columns = ['label_encoded'])
X = scaler.fit_transform(X)

y = df['label_encoded']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

# Correlation Matrix

label_encoded	1.000000	mfcc3_var	0.110916	mfcc17_mean	0.004921
spectral_bandwidth_mean	0.376621	mfcc7_var	0.107591	mfcc18_var	-0.039150
rolloff_mean	0.369515	rms_var	0.101467	mfcc17_var	-0.040047
spectral_centroid_mean	0.360175	mfcc5_mean	0.100811	mfcc16_var	-0.044354
chroma_stft_mean	0.330370	mfcc3_mean	0.089573	mfcc4_mean	-0.048384
mfcc1_mean	0.326771	mfcc20_mean	0.083224	mfcc11_var	-0.049489
spectral_centroid_var	0.281352	mfcc18_mean	0.082595	harmony_mean	-0.050513
rolloff_var	0.260298	mfcc10_mean	0.074355	perceptr_mean	-0.053023
zero_crossing_rate_mean	0.243590	mfcc8_mean	0.071334	mfcc19_var	-0.059141
zero_crossing_rate_var	0.215464	mfcc8_var	0.070706	mfcc15_var	-0.067919
spectral_bandwidth_var	0.210972	mfcc13_mean	0.066110	mfcc14_var	-0.068083
mfcc9_mean	0.207059	mfcc15_mean	0.065622	chroma_stft_var	-0.079683
rms_mean	0.205471	mfcc1_var	0.059184	mfcc13_var	-0.089063
mfcc7_mean	0.198996	mfcc5_var	0.051668	mfcc12_var	-0.090699
perceptr_var	0.170666	mfcc14_mean	0.051521	mfcc20_var	-0.099627
mfcc11_mean	0.152093	mfcc9_var	0.041014	mfcc2_mean	-0.348035
mfcc4_var	0.144013	mfcc6_mean	0.039947	length	NaN
mfcc6_var	0.131331	mfcc16_mean	0.026503		
mfcc2_var	0.127176	mfcc10_var	0.021395		
harmony_var	0.115452	mfcc19_mean	0.013143		
mfcc12_mean	0.111270	tempo	0.012369	Name: label_encoded, dtype: float64	

# Method

```
def modelling(model, title = "Default"):  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)  
    acc = accuracy_score(y_test, y_pred)  
    pre = precision_score(y_test, y_pred, average='macro')  
    rec = recall_score(y_test, y_pred, average='macro')  
  
    if title is not None: print('Accuracy', title, ':', acc)  
  
    return acc, pre, rec
```

# Method

```
# Naive Bayes
nb = GaussianNB()
modelling(nb, "Naive Bayes")

# KNN
knn = KNeighborsClassifier(n_neighbors=19)
modelling(knn, "KNN")

# Random Forest
rf = RandomForestClassifier(n_estimators=1000, max_depth=10, random_state=0)
modelling(rf, "Random Forest")

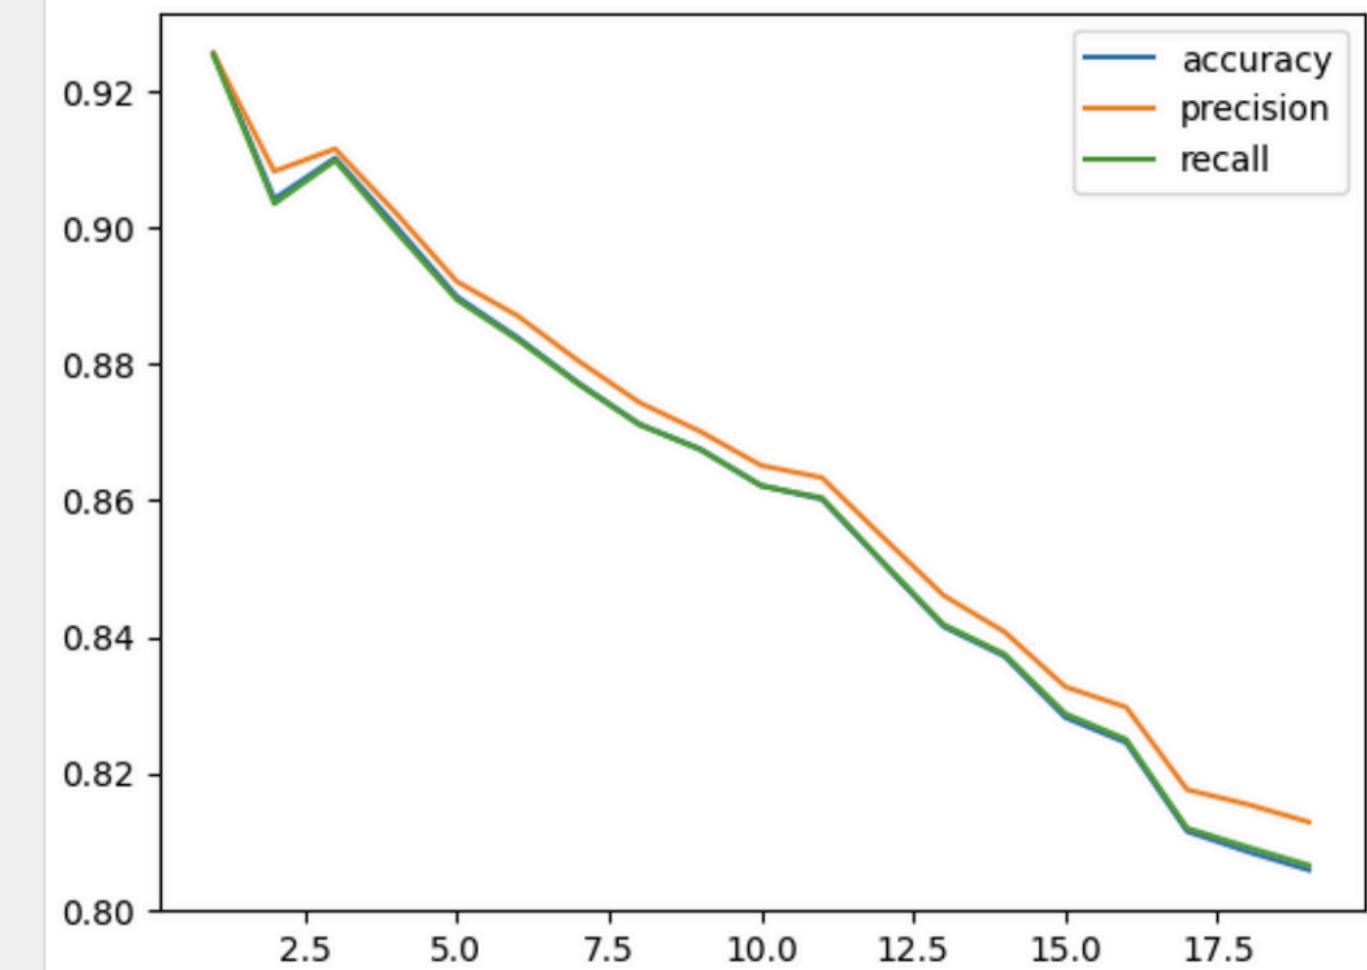
# Cross Gradient Booster
xgb = XGBClassifier(n_estimators=1000, learning_rate=0.05)
modelling(xgb, "Cross Gradient Booster")
```

# Model Evaluation KNN

```
acc = []
pre = []
rec = []

for i in range(1, 20):
    ac, pr, re = modelling(KNeighborsClassifier(n_neighbors=i),
None)
    acc.append(ac)
    pre.append(pr)
    rec.append(re)

x = range(1,20)
plt.plot(x, acc, label='accuracy')
plt.plot(x, pre, label='precision')
plt.plot(x, rec, label='recall')
plt.legend()
plt.show()
```



# Evaluation

Accuracy Naive Bayes : 0.5195195195195195

Accuracy KNN : 0.8058058058058059

Accuracy Random Forest : 0.8141474808141475

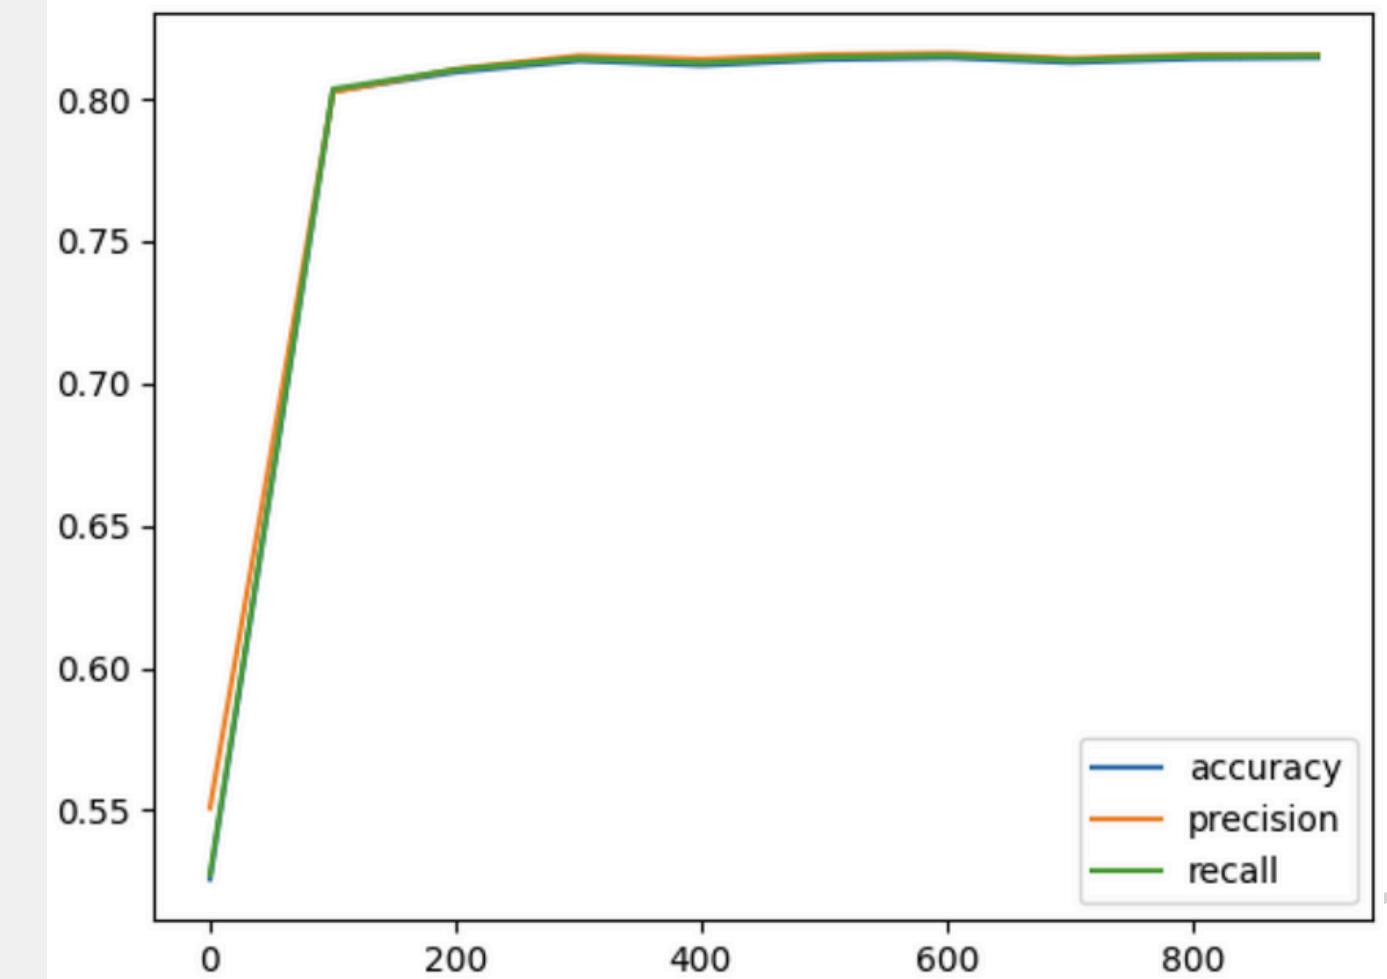
Accuracy Cross Gradient Booster : 0.9009009009009009

# Model Evaluation Random Forest

```
acc = []
pre = []
rec = []

for i in range(1, 1000, 100):
    ac, pr, re = modelling(RandomForestClassifier(n_estimators=i, max_depth=10, random_state=0), None)
    acc.append(ac)
    pre.append(pr)
    rec.append(re)

x = range(1, 1000, 100)
plt.plot(x, acc, label='accuracy')
plt.plot(x, pre, label='precision')
plt.plot(x, rec, label='recall')
plt.legend()
plt.show()
```

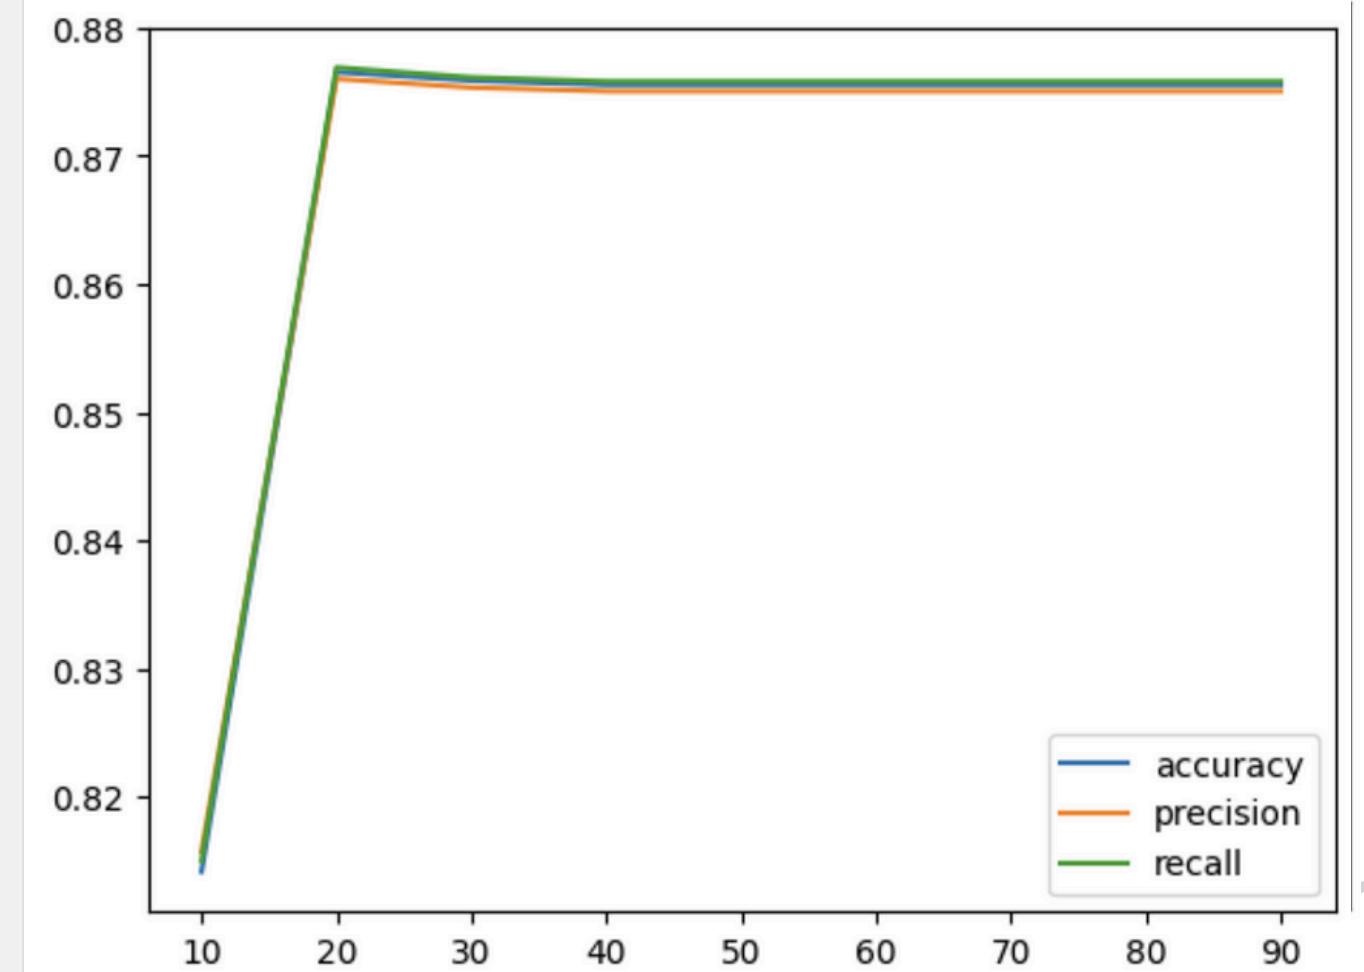


# Model Evaluation Random Forest

```
acc = []
pre = []
rec = []

for i in range(10, 100, 10):
    ac, pr, re = modelling(RandomForestClassifier(n_estimators=600, max_depth=i, random_state=0), None)
    acc.append(ac)
    pre.append(pr)
    rec.append(re)

x = range(10, 100, 10)
plt.plot(x, acc, label='accuracy')
plt.plot(x, pre, label='precision')
plt.plot(x, rec, label='recall')
plt.legend()
plt.show()
```

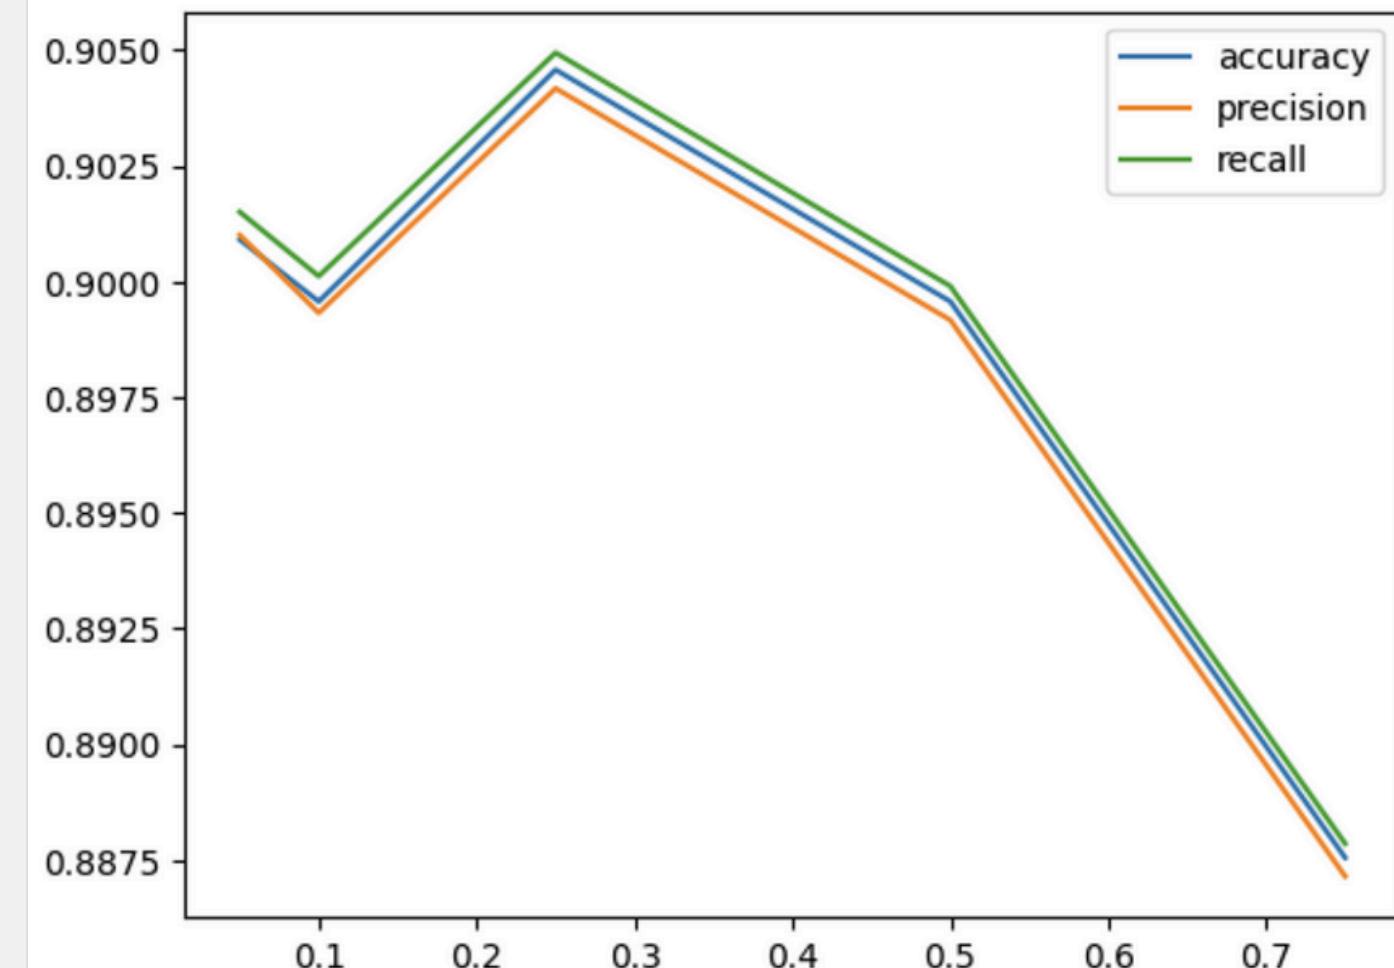


# Model Evaluation GB Classifier

```
acc = []
pre = []
rec = []

for i in [0.05, 0.1, 0.25, 0.5, 0.75]:
    ac, pr, re = modelling(XGBClassifier(n_estimators=1000, learning_rate=i), None)
    acc.append(ac)
    pre.append(pr)
    rec.append(re)

x = [0.05, 0.1, 0.25, 0.5, 0.75]
plt.plot(x, acc, label='accuracy')
plt.plot(x, pre, label='precision')
plt.plot(x, rec, label='recall')
plt.legend()
plt.show()
```

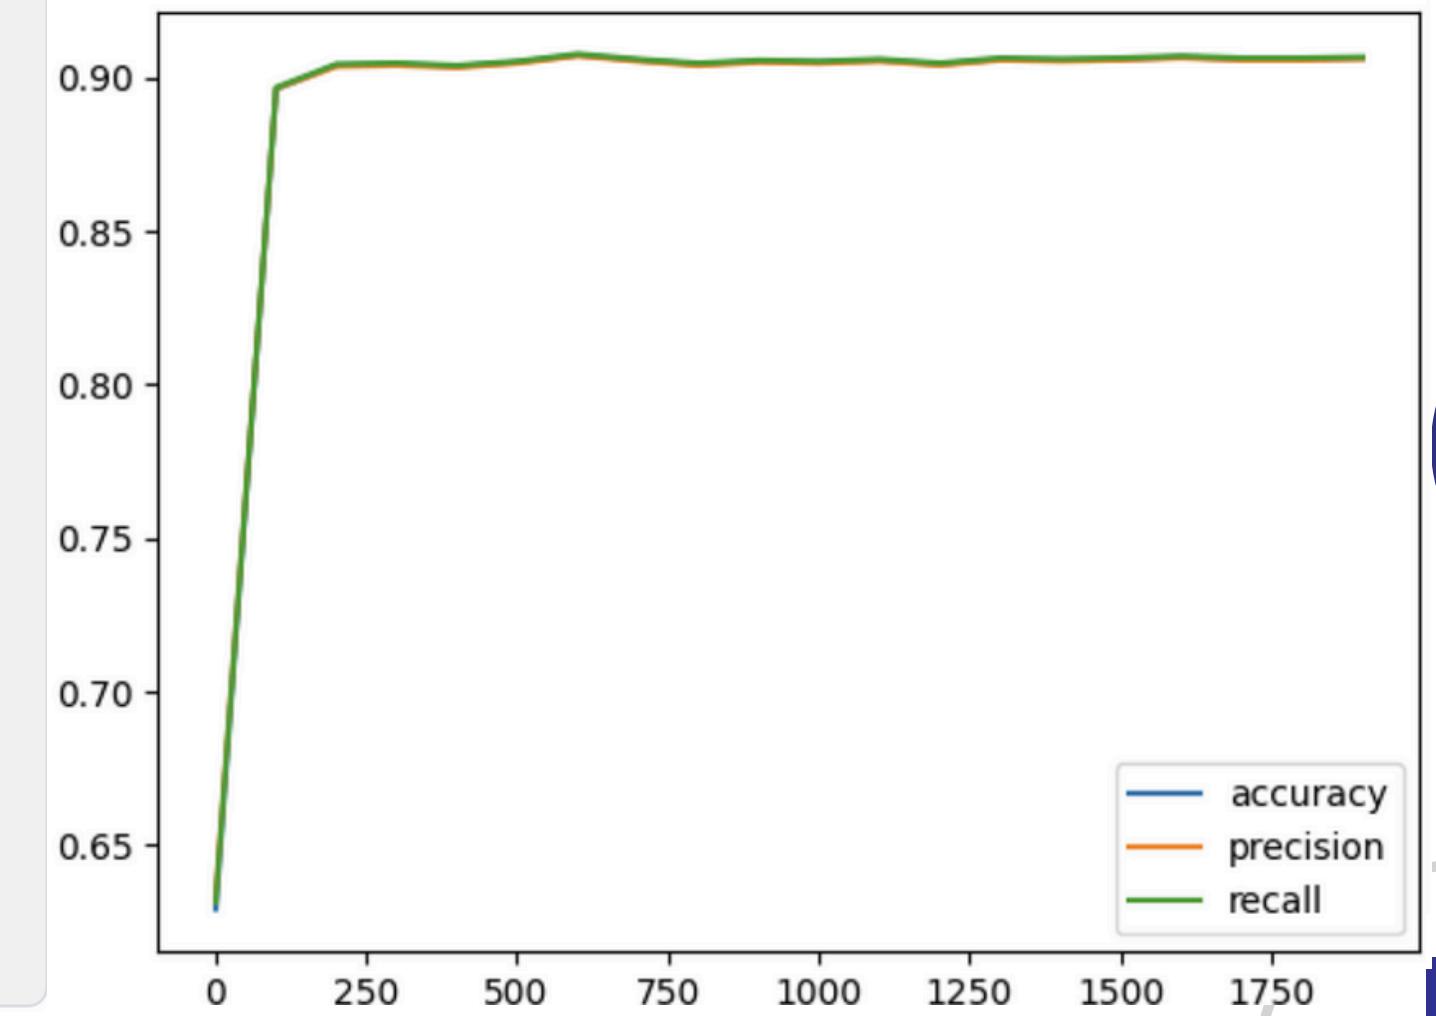


# Model Evaluation GB Classifier

```
acc = []
pre = []
rec = []

for i in range(1, 2000, 100):
    ac, pr, re = modelling(XGBClassifier(n_estimators=i, learning_rate=0.25), None)
    acc.append(ac)
    pre.append(pr)
    rec.append(re)

x = range(1, 2000, 100)
plt.plot(x, acc, label='accuracy')
plt.plot(x, pre, label='precision')
plt.plot(x, rec, label='recall')
plt.legend()
plt.show()
```



# Kesimpulan

- Setelah di-tuning peforma model menjadi
  - KNN ( $k = 1$ ) -> 92.5%
  - Random Forest ( $n\_estimator = 600$ ,  $max\_depth = 20$ ) -> 87.8%
  - XGBClassifier ( $n\_estimator = 1000$ ,  $learning\_rate = 0.25$ ) -> 90.5%

Saran:

Agar mendapatkan hasil model yang lebih baik lagi, kami menyarankan untuk menambah jumlah dari dataset, dan juga mengkombinasikan durasi lagu baik yang memiliki durasi yang panjang maupun yang pendek saat ingin melakukan feature extraction pada masa training, agar model dapat memprediksi lagu dengan durasi yang cukup panjang maupun pendek

# Demo

```
audio, sample_rate = librosa.load('/kaggle/input/gtzan-dataset-music-genre-classification/Data/genres_original/pop/pop.00050.wav')
print('Audio: ', audio)
print('Sample Rate: ', sample_rate)
print('Audio shape: ', audio.shape)

audio
ipd.display(ipd.Audio(audio, rate = sample_rate))
```

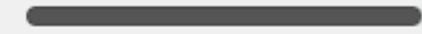
Audio: [0.03704834 0.03735352 0.037323 ... 0.05224609 0.0  
5944824 0.12362671]  
Sample Rate: 22050  
Audio shape: (661504,)

# Demo

```
ipd.display(ipd.Audio(audio, rate = sample_rate))
```



0:00 / 0:30



# Demo

```
# Harmonics and Perceptual
y_harmonic, y_perceptual = librosa.effects.hpss(audio)
print('Harmonics: ', y_harmonic)
print('Harmonics shape: ', y_harmonic.shape)
print('Harmonics mean: ', y_harmonic.mean())
print('Harmonics var: ', y_harmonic.var())
print('Perceptual: ', y_perceptual)
print('Perceptual shape: ', y_perceptual.shape)
print('Perceptual mean: ', y_perceptual.mean())
print('Perceptual var: ', y_perceptual.var())
```

```
Harmonics: [0.03326376 0.03181361 0.03337663 ... 0.01731614 0.00584044 0.02411156]
Harmonics shape: (661504,)
Harmonics mean: -3.5859378e-07
Harmonics var: 0.02375462
Perceptual: [0.00378458 0.00553991 0.00394636 ... 0.03492994 0.0536078 0.09951515]
Perceptual shape: (661504,)
Perceptual mean: -3.8225327e-05
Perceptual var: 0.0125090685
```

# Demo

```
# MFCC
mfcc = librosa.feature.mfcc(y = audio, sr = sample_rate)
print('MFCC: ', mfcc)
print('MFCC Shape: ', mfcc.shape)
print('MFCC Mean: ', mfcc.mean())
print('MFCC Var: ', mfcc.var())
plt.figure(figsize = (16, 6))
librosa.display.specshow(mfcc, sr=sample_rate, x_axis='time');
```

```
MFCC: [[ -232.74675   -125.48463    -56.231228 ...  -46.20056   -33.402805
         -31.713644 ]
[  48.240204    44.47583     53.990505 ...   79.02171    56.27352
   54.7958   ]
[  13.325691    22.516548    41.636967 ...  -30.524948   -13.2747345
   -1.5740416]
...
[ -12.231359   -6.256031     3.6848311 ...   -8.87232   -5.096953
   -3.8704803]
[ -9.302217    2.3334775     6.0767984 ...   -1.0246818   -4.763948
   -7.672885 ]
[  4.2557316   6.3851395     1.1116682 ...  -12.019724   -5.4266586
   -3.7294238]]
MFCC Shape: (20, 1293)
MFCC Mean: 4.024983
MFCC Var: 1010.4004
```

# Demo

```
[84]:  
def train_model(model):  
    model.fit(X_train, y_train)  
    return model  
  
knn = train_model(KNeighborsClassifier(n_neighbors=1))  
rf = train_model(RandomForestClassifier(n_estimators=600, max_depth=20, random_state=0))  
xgbc = train_model(XGBClassifier(n_estimators=600, learning_rate=0.25))  
  
+ Code + Markdown
```

# Demo

```
[92]: def extract_audio_features(file_path):
    y, sr = librosa.load(file_path, sr=None)

    features = {
        'length': y.shape[0] / sr,
        'chroma_stft_mean': np.mean(librosa.feature.chroma_stft(y=y, sr=sr)),
        'chroma_stft_var': np.var(librosa.feature.chroma_stft(y=y, sr=sr)),
        'rms_mean': np.mean(librosa.feature.rms(y=y)),
        'rms_var': np.var(librosa.feature.rms(y=y)),
        'spectral_centroid_mean': np.mean(librosa.feature.spectral_centroid(y=y, sr=sr)),
        'spectral_centroid_var': np.var(librosa.feature.spectral_centroid(y=y, sr=sr)),
        'spectral_bandwidth_mean': np.mean(librosa.feature.spectral_bandwidth(y=y, sr=sr)),
        'spectral_bandwidth_var': np.var(librosa.feature.spectral_bandwidth(y=y, sr=sr)),
        'rolloff_mean': np.mean(librosa.feature.spectral_rolloff(y=y, sr=sr)),
        'rolloff_var': np.var(librosa.feature.spectral_rolloff(y=y, sr=sr)),
        'zero_crossing_rate_mean': np.mean(librosa.feature.zero_crossing_rate(y)),
        'zero_crossing_rate_var': np.var(librosa.feature.zero_crossing_rate(y)),
        'harmony_mean': np.mean(librosa.effects.harmonic(y)),
        'harmony_var': np.var(librosa.effects.harmonic(y)),
        'percptr_mean': np.mean(librosa.effects.percussive(y)),
        'percptr_var': np.var(librosa.effects.percussive(y)),
        'tempo': librosa.beat.tempo(y=y, sr=sr)[0],
    }

    mfcc = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=20)
    for i in range(1, 21):
        features[f'mfcc{i}_mean'] = np.mean(mfcc[i-1])
        features[f'mfcc{i}_var'] = np.var(mfcc[i-1])

    return pd.DataFrame(features, index=[0])
```

# Demo

```
[93]:  
    file_path = '/kaggle/input/tes-lagu/Baby-Justin Bieber.mp3'  
    test_audio = extract_audio_features(file_path)  
  
    ▾ Expand  
  
[99]:  
    print(dict(zip(le.classes_, le.transform(le.classes_)))  
  
    {'blues': 0, 'classical': 1, 'country': 2, 'disco': 3, 'hiphop': 4, 'jazz': 5, 'metal': 6, 'pop': 7, 'reggae': 8, 'rock': 9}  
  
[98]:  
    col = test_audio.columns  
    test_audio = pd.DataFrame(scaler.transform(test_audio), columns=col)  
    print(knn.predict(test_audio))  
    print(rf.predict(test_audio))  
    print(xgbc.predict(test_audio))  
  
[4]  
[4]  
[8]
```

KNN: Hiphop, Random Forest: Hiphop, Xgbc: Reggae



# Terimakasih

